

---

# GeoLRM: Geometry-Aware Large Reconstruction Model for High-Quality 3D Gaussian Generation

---

Chubin Zhang<sup>1,2</sup> Hongliang Song<sup>2</sup> Yi Wei<sup>1</sup>

Yu Chen<sup>2</sup> Jiwen Lu<sup>1</sup> Yansong Tang<sup>1,†</sup>

<sup>1</sup>Tsinghua University

<sup>2</sup>Alibaba Group

{zhangcb19, y-wei19}@mails.tsinghua.edu.cn,

{hongliang.shl, chen.yu.chen}@alibaba-inc.com,

lujiwen@tsinghua.edu.cn, tang.yansong@sz.tsinghua.edu.cn.

<sup>†</sup> corresponding author

## Abstract

In this work, we introduce the Geometry-Aware Large Reconstruction Model (GeoLRM), an approach which can predict high-quality assets with 512k Gaussians and 21 input images in only 11 GB GPU memory. Previous works neglect the inherent sparsity of 3D structure and do not utilize explicit geometric relationships between 3D and 2D images. This limits these methods to a low-resolution representation and makes it difficult to scale up to the dense views for better quality. GeoLRM tackles these issues by incorporating a novel 3D-aware transformer structure that directly processes 3D points and uses deformable cross-attention mechanisms to effectively integrate image features into 3D representations. We implement this solution through a two-stage pipeline: initially, a lightweight proposal network generates a sparse set of 3D anchor points from the posed image inputs; subsequently, a specialized reconstruction transformer refines the geometry and retrieves textural details. Extensive experimental results demonstrate that GeoLRM significantly outperforms existing models, especially for dense view inputs. We also demonstrate the practical applicability of our model with 3D generation tasks, showcasing its versatility and potential for broader adoption in real-world applications. Our project page: GeoLRM Homepage.

## 1 Introduction

In fields ranging from robotics to virtual reality, the quality and diversity of 3D assets can dramatically influence both user experience and system efficiency. Historically, the creation of these assets has been a labour-intensive process, demanding the skills of expert artists and developers. While recent years have witnessed groundbreaking advancements in 2D image generation technologies, such as diffusion models [43, 44, 42] which iteratively refine images, their adaptation to 3D asset creation remains challenging. Directly applying diffusion models to 3D generation [20, 36] is less than satisfactory, primarily due to a dearth of large-scale and high-quality data. DreamFusion [40] innovatively optimize a 3D representation [2] by distilling the score of image distribution from pre-trained image diffusion models [43, 44]. However, this approach lacks a deep integration of 3D-specific knowledge, such as geometric consistency and spatial coherence, leading to significant issues such as the multi-head problem and the inconsistent 3D structure. Additionally, these methods require extensive per-scene optimizations, which severely limits their practical applications.

The introduction of the comprehensive 3D dataset Objaverse [12, 11] brings significant advancements for this field. Utilizing this dataset, researchers have fine-tuned 2D diffusion models to produce images

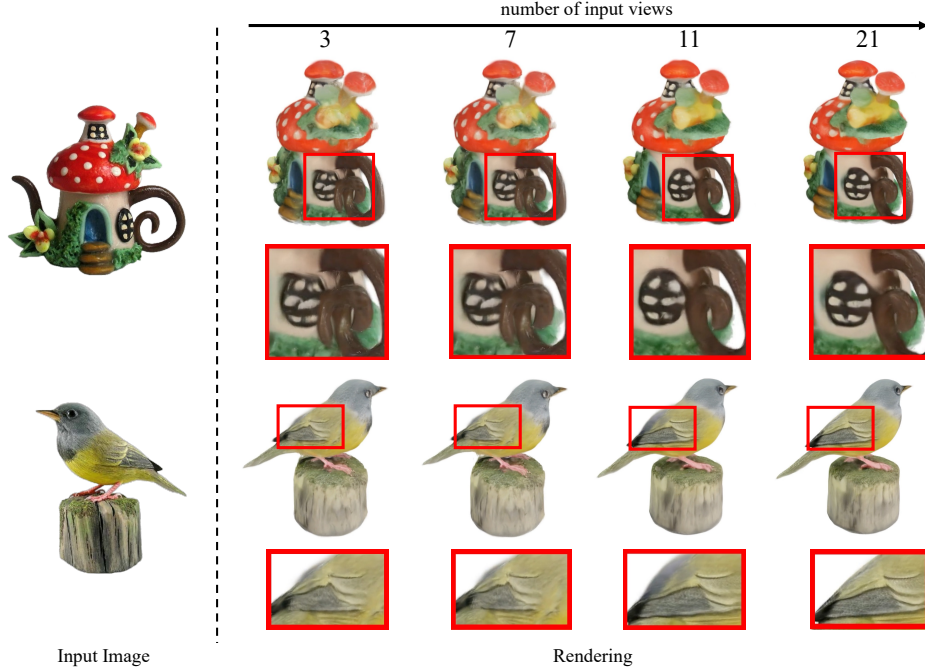


Figure 1: Image to 3D using GeoLRM. Initially, a 3D-aware diffusion model, specifically SV3D [60], transforms an input image into multiple views. Subsequently, these views are processed by our GeoLRM to generate detailed 3D assets. **Unlike other LRM-based approaches, GeoLRM notably improves as the number of input views increases.**

consistent with 3D structures [28, 47, 48]. Moreover, recent innovations [73, 64, 54, 71, 65] have combined these 3D-aware models with large reconstruction models (LRMs) [18] to achieve rapid and accurate 3D image generation. These methods typically employ large transformers or UNet models that convert sparse-view images into 3D representations in a single forward step. While they excel in speed and maintaining 3D consistency, they confront two primary limitations. Firstly, previous works utilize triplanes [18, 71, 64] to represent the 3D models, wasting lots of features in regions devoid of actual content and involving dense computations during rendering. This *violates the sparse nature of 3D* as our analysis shows that the visible portions of the 3D models in the Objaverse dataset constitute only about 5% of the overall spatial volume. Though Gaussian-based methods [54, 73, 65] may use pixel-aligned Gaussians for better efficiency, this representation is incapable of recovering the unseen area and thus heavily relies on the input images. Secondly, previous works tend to *overlook the explicit geometric relationships between 3D and 2D images*, which results in ineffective processing. The tri-plane or pixel-aligned Gaussian tokens do not correspond to a specific space in 3D, thus being unable to utilize the projection relationship between 3D points and images. In other words, they conduct dense attention between the 3D queries and the image keys. This leads to the fact that these methods tend to reconstruct 3D with sparse view inputs but cannot achieve better performance with denser inputs.

To address these challenges, we introduce the geometry-aware large reconstruction model (GeoLRM) for 3D Gaussian generation. Our method centres on a 3D-aware reconstruction transformer that eschews conventional representations like triplanes or pixel-aligned Gaussians in favour of a direct interaction within the 3D space. However, directly generating 3D Gaussians in the whole 3D space requires huge memory cost. To this end, we first propose a specialized proposal network to predict an occupancy grid from input images. Only the occupied voxels will be further processed to generate 3D Gaussian features. The proposed transformer replaces the dense cross attention with deformable cross attention [85]. By projecting the input 3D tokens onto the corresponding image planes, these tokens only focus on the most relevant features, which greatly improves the effectiveness.

We trained our GeoLRM on the Objaverse dataset rendered by [41] and tested it on the Google Scanned Objects [13]. By integrating geometric principles, our model not only outperforms existing methods with the same number of inputs but also makes it possible to work with denser image

inputs. Significantly, the model efficiently handles up to 21 images (even more if necessary), yielding superior 3D models in comparison to those generated from fewer images. Leveraging this capability, we integrated GeoLRM with SV3D [60] for high-quality 3D model generation.

## 2 Related Work

### 2.1 Optimization-based 3D reconstruction

3D reconstruction from multi-view images has been extensively studied in computer vision for decades. While traditional methods like SfM [68, 58, 45] and MVS [46, 16] provide basic reconstruction and calibration, they lack robustness and expressiveness. Recent advancements leverage learning-based methods for better performance. Among these methods, NeRF [33] stands out for its capability of capturing high-frequency details. Following works [2, 82, 3, 34, 76, 8, 53, 4] further improve its performance and speed. Though NeRF has made a great improvement, the need to query tons of points during the rendering process makes it hard for real-time applications. 3D Gaussians [21] solves this problem by explicitly expressing a scene with 3D Gaussians and utilizing an efficient rasterization pipeline. These methods involve a per-scene optimization process and require dense multi-view images for a good reconstruction.

### 2.2 Large Reconstruction Model

Different with optimization-based 3D reconstruction methods, large reconstruction models [18, 22, 54, 73, 65, 81, 62, 64] are able to reconstruct 3D shapes in a feed-forward way. As the pioneer work of this area, the LRM [18] illustrates that the transformer backbone can effectively leverage the power of large-scale datasets and translate image tokens into implicit 3D triplanes under multi-view supervision. Beyond LRM, Instant3D [22] improves reconstruction quality with sparse-view inputs. It employs a two-stage paradigm, which first generates four views with the diffusion model and then regresses NeRF [33] from generated multi-view images. Instead of NeRF, InstantMesh [71] utilizes mesh representation to reconstruct 3D objects, which adopts a differentiable iso-surface extraction module. However, many of works [54, 81, 73, 70] choose 3D Gaussians [21] as the outputs. GRM [73] proposes a transformer network to translate pixels to the set of pixel-aligned 3D Gaussians while LGM [54] uses an asymmetric UNet to predict and fuse 3D Gaussians. Compared with these methods, our GeoLRM projects multi-view features to the 3D space with cross-view attention mechanisms, which explicitly explores geometric knowledge.

### 2.3 3D generation

Early methods [6, 7, 15, 35, 51, 72, 37] in 3D generation area utilize 3D GANs to generate 3D-aware contents. Despite that some methods [32, 32, 84, 30, 10, 49, 79] replace 3D GANs with 3D diffusion models for high-quality generation, their generalization ability is bounded by the limited training data. Recently, proposed in DreamFusion [40], score distillation sampling (SDS) requires no 3D data and is able to leverage the great power of 2D text-to-image diffusion models [44, 43, 42]. Specifically, it optimizes a randomly-initialized 3D model and diffuses the render images with a pretrained diffusion model. As the follow-up works [63, 9, 26, 61, 55, 75, 27, 77, 25, 23, 41], many methods have been proposed to accelerate the optimization process or improve 3D generation quality. Different with SDS-based methods, Zero-1-to-3 [28] finetunes the 2D diffusion models on a large-scale synthetic dataset to change the camera viewpoint of a given image. Similar to Zero-1-to-3, many other works [47, 60, 48, 74, 29, 67, 31, 69] aim to synthesize multi-view consistent images. Our method can reconstruct 3D contents based on these synthesis multi-view images.

## 3 Methodology

### 3.1 Overview

Figure 2 illustrates the pipeline of our proposed method. Our approach takes a set of images  $\{I^i\}_{i=1}^N$  with their corresponding intrinsic  $\{K^i\}_{i=1}^N$  and extrinsic  $\{T^i\}_{i=1}^N$  as input. Initially, a proposal transformer predicts an occupancy grid. Each occupied voxel within this grid is considered a 3D anchor point. These 3D anchor points are then processed by a reconstruction transformer, refining

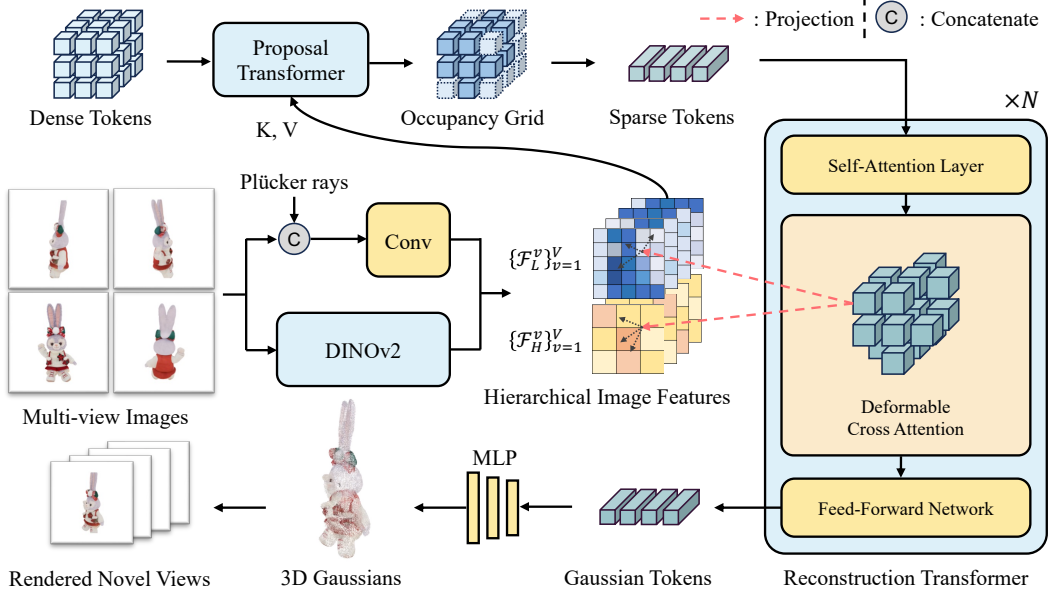


Figure 2: **Pipeline of the proposed GeoLRM**, a geometry-powered method for efficient images to 3D reconstruction. The process begins with the transformation of dense tokens into an occupancy grid via a Proposal Transformer, which captures spatial occupancy from hierarchical image features extracted using a combination of a convolutional layer and DINOv2 [38]. Sparse tokens representing occupied voxels are further processed through a Reconstruction Transformer that employs self-attention and deformable cross-attention mechanisms to refine geometry and retrieve texture details with 3D to 2D projection. Finally, the refined 3D tokens are converted into 3D Gaussians for real-time rendering.

their geometry and retrieving textural details. The proposal and reconstruction transformers share the same model architecture, which is further discussed in Section 3.2. The outputs of the reconstruction transformer are decoded into Gaussian features with a shallow MLP for rendering. Loss functions are described in Section 3.3.

### 3.2 Model Architecture

We present a geometry-aware transformer architecture featuring a hierarchical image encoder for extracting high and low-level image feature maps, and an anchor point decoder for transforming these features into 3D representations.

**Hierarchical Image Encoder** Our method integrates both high and low-level features to enhance model performance. For high-level features, we utilize DINOv2 [38], which excels in single-image 3D tasks [1]. To capture low-level features, we combine Plücker ray embeddings and RGB values. The Plücker ray parameterizes each ray corresponding to a pixel by  $\mathbf{r} = (\mathbf{d}, \mathbf{o} \times \mathbf{d})$ , with  $\mathbf{d}$  representing the ray’s direction and  $\mathbf{o}$  its origin [50, 74]. These embeddings, denoted as  $R^v$  for each image  $I^v$ , are concatenated with the RGB values of the image. This combined data is then integrated through a convolution layer. The encoding processes are succinctly described by the equations:

$$\mathcal{F}_H^v = \text{DINOv2}(I^v), \quad (1)$$

$$\mathcal{F}_L^v = \text{Conv}(\text{Concat}(I^v, R^v)), \quad (2)$$

where  $\mathcal{F}_H^v$  and  $\mathcal{F}_L^v$  represent the high and low-level feature maps of image  $I^v$ , respectively.

**Anchor Point Decoder** The anchor point decoder aims to efficiently lift image features to 3D. Previous methods [18, 54, 73, 65, 81] uses tri-planes or pixel-aligned Gaussians to represent 3D contents. However, these data structures make it hard to utilize the projection relationships, causing dense computations. Instead, we use 3D anchor points, which serve as proxies for their surrounding points, significantly reducing the number of points we need to process. As detailed in Figure 2, each decoder block contains a self-attention layer, a deformable cross-attention layer and a feed-forward

network (FFN). The model takes  $N$  anchor point features  $\mathcal{F}_A = \{\mathbf{f}_i\}_{i=1}^N$  as input tokens. Each token  $\mathbf{f}_i$  comprises the coordinate of the corresponding point and a shared learnable feature.

For the self-attention layer, a crucial problem is how to inject positional information into the sparse 3D tokens. We extend the Rotary Positional Embedding (RoPE) [52] to 3D conditions for relative positional embedding. For a query  $\mathbf{q}_m$  and a key  $\mathbf{k}_n$  at absolute position  $m$  and  $n$ , we ensure that the inner product of embedded values reflects only the relative position information  $m - n$ . A direct yet promising way is splitting the features into three parts and applying RoPE [52] on each part with x, y, and z positions respectively.

As we can locate each anchor point in the 3D space, a possible way to lift 2D features to 3D is to project them to the feature maps with known poses and average the corresponding features. However, this method assumes an accurate anchor position, an equal contribution of all images and a good 3D correspondence of input images, which is often impractical, especially in 3D generation tasks. To tackle these issues, we employ deformable attention [85, 24, 66] for a robust fusion of image features. Given a 3D anchor point feature  $\mathbf{f}_i$ , its spatial coordinate  $\mathbf{x}_i$  and multiple feature maps  $\{\mathcal{F}^v\}_{v=1}^V$ , the deformable attention mechanism is formulated as:

$$\text{DeformAttn}(\mathbf{f}_i, \mathbf{x}_i, \{\mathcal{F}^v\}_{v=1}^V) = \sum_{v=1}^V w_v \left[ \sum_{k=1}^K A_k \mathcal{F}^v \langle \mathbf{p}_{iv} + \Delta \mathbf{p}_{ivk} \rangle \right], \quad (3)$$

where  $k$  indexes the sampled keys and  $K$  is the total sampled key numbers.  $\mathbf{p}_{iv}$  is the projected 2D coordinate on feature map  $\mathcal{F}^v$  and  $\Delta \mathbf{p}_{ivk}$  is the sampled offset.  $\langle \cdot \rangle$  indicates the interpolation operation.  $A_k$  is the attention weight predicted from  $\mathbf{f}_i$ .  $w_v$  is a per-view weight derived from the feature it weights. Notably, the prediction of  $\Delta \mathbf{p}_{ivk}$  allows the network to correct the geometry error of anchor points and the inconsistency of input images; The  $w_v$  enables different importance levels for each image. To further enhance the representation ability of the model, this mechanism is extended to multi-head and multi-scale conditions.

Given input tokens  $\mathcal{F}_A^{in}$ , the decoder block enhances these tokens through a series of sophisticated transformations:

$$\mathcal{F}_A^{self} = \mathcal{F}_A^{in} + \text{SelfAttn}(\text{RMSNorm}(\mathcal{F}_A^{in})), \quad (4)$$

$$\mathcal{F}_A^{cross} = \mathcal{F}_A^{self} + \text{DeformCrossAttn}(\text{RMSNorm}(\mathcal{F}_A^{self}), \{(\mathcal{F}_H^v, \mathcal{F}_L^v)\}_{v=1}^V), \quad (5)$$

$$\mathcal{F}_A^{out} = \mathcal{F}_A^{cross} + \text{FFN}(\text{RMSNorm}(\mathcal{F}_A^{cross})). \quad (6)$$

This design introduces several improvements over the original transformer architecture [59]. By incorporating RMSNorm [78] for normalization and SiLU [14] for activation, we achieve more stable training dynamics and better performance.

**Post Processing** The proposal network takes a low-resolution dense grid ( $16^3$ ) as anchor points. The output is upsampled to a high-resolution grid ( $128^3$ ) with a linear layer. This grid is formulated to represent the occupancy probability of the corresponding area ( $[-0.5, 0.5]^3$ ). The reconstruction transformer takes occupied voxels as anchor points. Each output token  $\mathbf{f}_i$  is decoded into multiple 3D Gaussians  $\{\mathbf{G}_{ij}\}_{j=1}^M$  with a linear layer. The 3D Gaussian  $\mathbf{G}_{ij}$  is parameterized by the offset  $\mathbf{o}_{ij}$  regarding the anchor points, 3-channel RGB  $\mathbf{c}_{ij}$ , 3-channel scale  $\mathbf{s}_{ij}$ , 4-channel rotation quaternion  $\boldsymbol{\sigma}_{ij}$ , and 1-channel opacity  $\alpha_{ij}$ . We employ activation functions to limit the range of the offset, scale and opacity for better training stability similar to [54]:

$$\mathbf{o}_{ij} = \text{Sigmoid}(\mathbf{o}'_{ij}) \cdot \mathbf{o}_{\max}, \quad (7)$$

$$\mathbf{s}_{ij} = \text{Sigmoid}(\mathbf{s}'_{ij}) \cdot \mathbf{s}_{\max}, \quad (8)$$

$$\alpha_{ij} = \text{Sigmoid}(\alpha'_{ij}), \quad (9)$$

where  $\mathbf{o}_{\max}$ ,  $\mathbf{s}_{\max}$  are predefined maximum values of offsets and scales. Given target camera views  $\{\mathbf{c}_t\}_{t=1}^T$ , the 3D Gaussians can be further rendered into images  $\{\hat{\mathbf{I}}_t\}_{t=1}^T$ , alpha masks  $\{\hat{\mathbf{M}}_t\}_{t=1}^T$  and depth maps  $\{\hat{\mathbf{D}}_t\}_{t=1}^T$  through Gaussian splatting [21].

### 3.3 Training Objectives

We employ a two-stage training mechanism for our model. In the first stage, we train the proposal transformer using 3D occupancy ground truth. This stage presents a challenge as it involves a highly

unbalanced binary classification task; only about 5% of the voxels are occupied. To address this imbalance, we employ a combination of binary cross-entropy loss and the scene-class affinity loss, as proposed in [5], to supervise the training process. For the generation of ground truth data, see A.2.

For the second stage, we supervise the rendered  $T$  images, alpha masks and depth maps with corresponding ground truth:

$$\mathcal{L} = \sum_{t=1}^T \left( \mathcal{L}_{\text{img}}(\hat{I}_t, I_t) + \mathcal{L}_{\text{mask}}(\hat{M}_t, M_t) + 0.2\mathcal{L}_{\text{depth}}(\hat{D}_t, D_t, I_t) \right), \quad (10)$$

$$\mathcal{L}_{\text{img}}(\hat{I}_t, I_t) = \|\hat{I}_t - I_t\|_2 + 2\mathcal{L}_{\text{LPIPS}}(\hat{I}_t, I_t), \quad (11)$$

$$\mathcal{L}_{\text{mask}}(\hat{M}_t, M_t) = \|\hat{M}_t - M_t\|_2, \quad (12)$$

$$\mathcal{L}_{\text{depth}}(\hat{D}_t, D_t, I_t) = \frac{1}{|\hat{D}_t|} \left\| \exp(-\Delta I_t) \odot \log(1 + |\hat{D}_t - D_t|) \right\|_1, \quad (13)$$

where  $\mathcal{L}_{\text{LPIPS}}$  is the perceptual image patch similarity loss [83],  $|\hat{D}_t|$  is the total number of pixels in  $|\hat{D}_t|$ ,  $\Delta I_t$  is the gradient of the current RGB image and  $\odot$  is the element-wise multiplication operation. As demonstrated in [57], applying a logarithmic penalty and weighting the per-pixel depth errors with the image gradients result in a smoother geometric representation.

## 4 Experiments

### 4.1 Datasets

**G-buffer Objaverse (GObjaverse) [41]:** Used for training. Derived from the original Objaverse [12] dataset, GObjaverse includes high-quality renderings of albedo, RGB, depth, and normal images. These images are generated through a hybrid technique combining rasterization and path tracing. The dataset comprises approximately 280,000 normalized 3D models scaled to fit within a cubic space of  $[-0.5, 0.5]^3$ . GObjaverse employs a diverse camera setup involving:

- Two orbital paths yielding 36 views per model. This includes 24 views at elevations between  $5^\circ$  and  $30^\circ$  (incremented by  $15^\circ$  rotations) and 12 views at near-horizontal elevations from  $-5^\circ$  to  $5^\circ$  (with  $30^\circ$  rotation steps).
- Additional top and bottom views for comprehensive spatial coverage.

**Google Scanned Objects (GSO) [13]:** Used for evaluation, this dataset is rendered similarly to GObjaverse to maintain consistency. We randomly select a subset of 100 objects to streamline the evaluation process.

### 4.2 Implimentation details

Our model features 330 million parameters distributed across two distinct image encoders and two transformers. The first encoder processes geometry with the 6-layer proposal transformer, while the second focuses more on textures crucial with the 16-layer reconstruction transformer. During training, we maintain a maximum number of transformer input tokens of 4k and randomly select 8 views from a possible 38 for supervision. From these 8 views, we randomly select 1 to 7 views as inputs to predict the remaining views. This flexibility in view selection not only tests the robustness of our method but also mimics real-world scenarios where complete data may not always be available. Both input and rendering resolutions are maintained at 448x448 pixels. At the testing and inference stages, the model processes up to 16k input tokens, showcasing its scalability without the need for fine-tuning. Detailed information on our model’s architecture and training procedures can be found in Section A.3.

### 4.3 Quantitative Results

We evaluated the quality of reconstructed assets from sparse view inputs by analyzing both 2D visual and 3D geometric aspects on the GSO dataset [13]. Visual quality was assessed by comparing rendered views to ground truth images using metrics such as PSNR, SSIM, and LPIPS. Geometric

Table 1: Quantitative results on Google Scanned Objects (GSO) [13], where we used six views for inputs and four for evaluation. **Bold** and underline denote the highest and second-highest scores, respectively.

Method	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	CD $\downarrow$	FS $\uparrow$
LGM [54]	20.76	0.832	0.227	0.295	0.703
CRM [64]	22.78	0.843	0.190	0.213	0.831
InstantMesh [71]	<u>23.19</u>	<u>0.856</u>	<b>0.166</b>	<u>0.186</u>	<u>0.854</u>
Ours	<b>23.42</b>	<b>0.865</b>	<u>0.174</u>	<b>0.165</b>	<b>0.890</b>

Table 2: Quantitative results on Google Scanned Objects (GSO) [13] with different numbers of input views. We keep the same four views for testing while changing the number of input views. **Bold** denotes the highest score.

Method	4 Inputs		8 Inputs		12 Inputs	
	PSNR $\uparrow$	SSIM $\uparrow$	PSNR $\uparrow$	SSIM $\uparrow$	PSNR $\uparrow$	SSIM $\uparrow$
InstantMesh [71]	<b>22.87</b>	0.832	23.22	0.861	23.05	0.843
Ours	22.64	<b>0.840</b>	<b>23.76</b>	<b>0.870</b>	<b>24.39</b>	<b>0.887</b>

accuracy was evaluated by aligning our models to the ground truth coordinate systems and measuring discrepancies using Chamfer Distance and F-Score at a threshold of 0.2, with point samples totalling 16,000 from the ground truth surfaces. Our method was quantitatively compared against established baselines, including LGM [54], CRM [64], and InstantMesh [71]. We avoided comparisons with proprietary methods due to the unavailability of their test splits. Similarly, we excluded comparisons with OpenLRM [17] and TripoSR [56] as these methods are tailored for single image inputs, which would be unfair to compare with.

Our approach achieved state-of-the-art performance in four out of the five metrics studied. Although InstantMesh showed slightly higher LPIPS, attributed to its mesh-based smoothing capabilities, our method demonstrated superior geometric accuracy, benefiting from explicit modelling of the 3D-to-2D relationship.

In another experiment, outlined in Table 2, we observed a notable trend: the performance of our model improves consistently as the number of input views increases. This indicates robust scalability, a critical feature for practical applications. In contrast, the performance of InstantMesh [71], does not follow this pattern. Specifically, InstantMesh shows a decline in performance when the input views increase to 12. This degradation could be due to two primary factors. First, the low-resolution triplanes may reach their maximum capacity to represent details. Second, the model tends to oversmooth details when handling a large volume of image tokens. Our approach strategically addresses these issues. We employ an extendable sequence of 3D tokens that can be dynamically adjusted to fit the resolution requirements. Additionally, our model features deformable attention mechanisms that intelligently focus on the most pertinent information, preventing the loss of critical details.

#### 4.4 Qualitative Results

We conducted a qualitative analysis comparing our method with several LRM-based baselines, including TripoSR [17], LGM [54], CRM [64], and InstantMesh [71], maintaining their original settings to ensure optimal performance. In our approach, we utilized the SV3D [60] technology to generate 21 multi-view images, significantly enhancing the resolution and textural details of the 3D Gaussians produced, as illustrated in Figure 3. Furthermore, as shown in Figure 4, employing InstantMesh to reconstruct these images did not yield satisfactory outcomes, corroborating our quantitative findings. This demonstrates the superior capability of our method in handling complex 3D reconstructions.



Figure 3: Qualitative comparisons of different image-3D methods. **Better viewed when zoomed in.**




Figure 4: Qualitative comparison concerning scalability in input views.

#### 4.5 Ablation Study

We provide ablation studies for the key designs of our method as shown in Table 3. Due to the limited computation sources, the ablation is done using a smaller reconstruction model (12 layers) and lower resolution (224x224).

**Hierarchical Image Encoder** Our ablation study underscores the critical role of hierarchical image features in reconstruction tasks, which necessitate both high-level semantic information (e.g., object identity and arrangement) and low-level texture information (e.g., surface patterns and colors). As illustrated in Figure 5, the absence of high-level features leads to model instability, while omitting



Table 3: Ablation study. Models are tested on the GSO dataset [13]. Upper: 6 input views and 4 testing views. Lower: different input views. **Bold** and underline denote the highest and second-highest scores, respectively.

Method	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$
W/o low-level features	20.29	0.817	0.246
W/o high-level features	15.85	0.798	0.289
W/o 3D RoPE	20.52	0.827	0.224
Fixed # input views	<b>20.97</b>	<b>0.839</b>	<u>0.220</u>
Full model	<u>20.73</u>	<u>0.831</u>	<b>0.216</b>

Method	4 Inputs		8 Inputs		12 Inputs	
	PSNR $\uparrow$	SSIM $\uparrow$	PSNR $\uparrow$	SSIM $\uparrow$	PSNR $\uparrow$	SSIM $\uparrow$
Fixed # input views	<u>19.72</u>	<u>0.822</u>	<u>20.85</u>	<u>0.833</u>	<u>21.43</u>	<u>0.838</u>
Full model	<b>19.94</b>	<b>0.835</b>	<b>21.16</b>	<b>0.840</b>	<b>22.04</b>	<b>0.853</b>




Figure 5: Effects of excluding high-level and low-level features in the image encoder.

low-level features results in a loss of textural detail. This dual requirement emphasizes the model’s reliance on a comprehensive feature set for accurate image reconstruction.

**3D RoPE** In transformer-based architectures, the role of positional embeddings is critical for accurately interpreting sequence data positions. A key question arises: With the reconstruction transformer employing deformable cross attention to elevate 2D features to 3D, is positional embedding still necessary? Our ablation studies confirm its necessity. Notably, 3D RoPE significantly enhances the model’s ability to handle longer sequences. For instance, increasing the sequence length from 4k to 16k elements, models equipped with 3D RoPE exhibited a PSNR improvement of 0.4, compared to a 0.2 improvement in models lacking 3D RoPE. This observation aligns with the 1D RoPE [52].

**Dynamic Input** The ablation study demonstrates a decrease in performance when employing our dynamic input view strategy compared to the fixed 6 input view setting when the training and testing phases were consistent. Despite this, the dynamic input strategy enhances the model’s ability to generalize across different input configurations. This adaptability is critical for handling more complex scenarios, aligning with our primary objectives.

## 5 Limitations

Although our two-stage method achieves high-quality reconstruction, it is not an end-to-end model, which leads to error accumulation. In other words, the results of the second stage highly depend on the occupancy grids of the first stage. Currently, we have to use the proposal network since directly processing Gaussian points in the whole 3D space is time-consuming. For the future work, we plan to extend our method in an end-to-end manner.

## 6 Conclusion

In this paper, we propose geometry-aware large reconstruction model for efficient and high-quality 3D generation. Different with previous works, our method explores the sparsity of 3D and leverages the explicit geometric relationship between 3D and 2D images. To this end, GeoLRM adopts a 3D-aware transformer structure to predict 3D Gaussians in a coarse-to-fine fashion. Specifically, we first utilize a proposal network to predict the coarse occupancy grids of 3D assets, which provide the initial 3D anchor points for the second stage. Then we employ the deformable cross attention to refine the 3D structure. Experimental results show that the proposed method can efficiently process higher resolution and denser image inputs with better performance.

## References

- [1] Mohamed El Banani, Amit Raj, Kevis-Kokitsi Maninis, Abhishek Kar, Yuanzhen Li, Michael Rubinstein, Deqing Sun, Leonidas Guibas, Justin Johnson, and Varun Jampani. Probing the 3d awareness of visual foundation models. *arXiv preprint arXiv:2404.08636*, 2024.
- [2] Jonathan T Barron, Ben Mildenhall, Matthew Tancik, Peter Hedman, Ricardo Martin-Brualla, and Pratul P Srinivasan. Mip-nerf: A multiscale representation for anti-aliasing neural radiance fields. In *ICCV*, pages 5855–5864, 2021.
- [3] Jonathan T Barron, Ben Mildenhall, Dor Verbin, Pratul P Srinivasan, and Peter Hedman. Mip-nerf 360: Unbounded anti-aliased neural radiance fields. In *CVPR*, pages 5470–5479, 2022.
- [4] Jonathan T Barron, Ben Mildenhall, Dor Verbin, Pratul P Srinivasan, and Peter Hedman. Zip-nerf: Anti-aliased grid-based neural radiance fields. In *ICCV*, pages 19697–19705, 2023.
- [5] Anh-Quan Cao and Raoul De Charette. Monoscene: Monocular 3d semantic scene completion. In *CVPR*, pages 3991–4001, 2022.
- [6] Eric R Chan, Connor Z Lin, Matthew A Chan, Koki Nagano, Boxiao Pan, Shalini De Mello, Orazio Gallo, Leonidas J Guibas, Jonathan Tremblay, Sameh Khamis, et al. Efficient geometry-aware 3d generative adversarial networks. In *CVPR*, pages 16123–16133, 2022.
- [7] Eric R Chan, Marco Monteiro, Petr Kellnhofer, Jiajun Wu, and Gordon Wetzstein. pi-gan: Periodic implicit generative adversarial networks for 3d-aware image synthesis. In *CVPR*, pages 5799–5809, 2021.
- [8] Anpei Chen, Zexiang Xu, Andreas Geiger, Jingyi Yu, and Hao Su. Tensorf: Tensorial radiance fields. In *ECCV*, pages 333–350. Springer, 2022.
- [9] Rui Chen, Yongwei Chen, Ningxin Jiao, and Kui Jia. Fantasia3d: Disentangling geometry and appearance for high-quality text-to-3d content creation. In *ICCV*, pages 22246–22256, 2023.
- [10] Gene Chou, Yuval Bahat, and Felix Heide. Diffusion-sdf: Conditional generative modeling of signed distance functions. In *ICCV*, pages 2262–2272, 2023.
- [11] Matt Deitke, Ruoshi Liu, Matthew Wallingford, Huong Ngo, Oscar Michel, Aditya Kusupati, Alan Fan, Christian Laforte, Vikram Voleti, Samir Yitzhak Gadre, et al. Objaverse-xl: A universe of 10m+ 3d objects. *NeurIPS*, 36, 2024.
- [12] Matt Deitke, Dustin Schwenk, Jordi Salvador, Luca Weihs, Oscar Michel, Eli VanderBilt, Ludwig Schmidt, Kiana Ehsani, Aniruddha Kembhavi, and Ali Farhadi. Objaverse: A universe of annotated 3d objects. In *CVPR*, pages 13142–13153, 2023.
- [13] Laura Downs, Anthony Francis, Nate Koenig, Brandon Kinman, Ryan Hickman, Krista Reymann, Thomas B McHugh, and Vincent Vanhoucke. Google scanned objects: A high-quality dataset of 3d scanned household items. In *ICRA*, pages 2553–2560. IEEE, 2022.
- [14] Stefan Elfving, Eiji Uchibe, and Kenji Doya. Sigmoid-weighted linear units for neural network function approximation in reinforcement learning. *Neural networks*, 107:3–11, 2018.
- [15] Jun Gao, Tianchang Shen, Zian Wang, Wenzheng Chen, Kangxue Yin, Daiqing Li, Or Litany, Zan Gojcic, and Sanja Fidler. Get3d: A generative model of high quality 3d textured shapes learned from images. *NeurIPS*, 35:31841–31854, 2022.
- [16] Michael Goesele, Brian Curless, and Steven M Seitz. Multi-view stereo revisited. In *CVPR*, volume 2, pages 2402–2409. IEEE, 2006.

- [17] Zexin He and Tengfei Wang. Openlrm: Open-source large reconstruction models. <https://github.com/3DTopia/OpenLRM>, 2023.
- [18] Yicong Hong, Kai Zhang, Jiuxiang Gu, Sai Bi, Yang Zhou, Difan Liu, Feng Liu, Kalyan Sunkavalli, Trung Bui, and Hao Tan. Lrm: Large reconstruction model for single image to 3d. In *ICLR*, 2023.
- [19] Yuanhui Huang, Wenzhao Zheng, Borui Zhang, Jie Zhou, and Jiwen Lu. Selfocc: Self-supervised vision-based 3d occupancy prediction. *arXiv preprint arXiv:2311.12754*, 2023.
- [20] Heewoo Jun and Alex Nichol. Shap-e: Generating conditional 3d implicit functions. *arXiv preprint arXiv:2305.02463*, 2023.
- [21] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *TOG*, 42(4):1–14, 2023.
- [22] Jiahao Li, Hao Tan, Kai Zhang, Zexiang Xu, Fujun Luan, Yinghao Xu, Yicong Hong, Kalyan Sunkavalli, Greg Shakhnarovich, and Sai Bi. Instant3d: Fast text-to-3d with sparse-view generation and large reconstruction model. In *ICLR*, 2023.
- [23] Weiyu Li, Rui Chen, Xuelin Chen, and Ping Tan. Sweetdreamer: Aligning geometric priors in 2d diffusion for consistent text-to-3d. In *CVPR*, 2024.
- [24] Zhiqi Li, Wenhai Wang, Hongyang Li, Enze Xie, Chonghao Sima, Tong Lu, Yu Qiao, and Jifeng Dai. Bevformer: Learning bird’s-eye-view representation from multi-camera images via spatiotemporal transformers. In *ECCV*, pages 1–18. Springer, 2022.
- [25] Yixun Liang, Xin Yang, Jiantao Lin, Haodong Li, Xiaogang Xu, and Yingcong Chen. Luciddreamer: Towards high-fidelity text-to-3d generation via interval score matching. In *CVPR*, 2024.
- [26] Chen-Hsuan Lin, Jun Gao, Luming Tang, Towaki Takikawa, Xiaohui Zeng, Xun Huang, Karsten Kreis, Sanja Fidler, Ming-Yu Liu, and Tsung-Yi Lin. Magic3d: High-resolution text-to-3d content creation. In *CVPR*, pages 300–309, 2023.
- [27] Fangfu Liu, Diankun Wu, Yi Wei, Yongming Rao, and Yueqi Duan. Sherpa3d: Boosting high-fidelity text-to-3d generation via coarse 3d prior. In *CVPR*, 2024.
- [28] Ruoshi Liu, Rundi Wu, Basile Van Hoorick, Pavel Tokmakov, Sergey Zakharov, and Carl Vondrick. Zero-1-to-3: Zero-shot one image to 3d object. In *ICCV*, pages 9298–9309, 2023.
- [29] Yuan Liu, Cheng Lin, Zijiao Zeng, Xiaoxiao Long, Lingjie Liu, Taku Komura, and Wenping Wang. Syncdreamer: Generating multiview-consistent images from a single-view image. In *ICLR*, 2024.
- [30] Zhen Liu, Yao Feng, Michael J Black, Derek Nowrouzezahrai, Liam Paull, and Weiyang Liu. Meshdiffusion: Score-based generative 3d mesh modeling. In *ICLR*, 2023.
- [31] Xiaoxiao Long, Yuan-Chen Guo, Cheng Lin, Yuan Liu, Zhiyang Dou, Lingjie Liu, Yuexin Ma, Song-Hai Zhang, Marc Habermann, Christian Theobalt, et al. Wonder3d: Single image to 3d using cross-domain diffusion. *arXiv preprint arXiv:2310.15008*, 2023.
- [32] Luke Melas-Kyriazi, Christian Rupprecht, and Andrea Vedaldi. Pc2: Projection-conditioned point cloud diffusion for single-image 3d reconstruction. In *CVPR*, pages 12923–12932, 2023.
- [33] Ben Mildenhall, Pratul P Srinivasan, Matthew Tancik, Jonathan T Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. *Communications of the ACM*, 65(1):99–106, 2021.
- [34] Thomas Müller, Alex Evans, Christoph Schied, and Alexander Keller. Instant neural graphics primitives with a multiresolution hash encoding. *TOG*, 41(4):102:1–102:15, July 2022.
- [35] Thu Nguyen-Phuoc, Chuan Li, Lucas Theis, Christian Richardt, and Yong-Liang Yang. Hologan: Unsupervised learning of 3d representations from natural images. In *ICCV*, pages 7588–7597, 2019.
- [36] Alex Nichol, Heewoo Jun, Prafulla Dhariwal, Pamela Mishkin, and Mark Chen. Point-e: A system for generating 3d point clouds from complex prompts. *arXiv preprint arXiv:2212.08751*, 2022.
- [37] Michael Niemeyer and Andreas Geiger. Giraffe: Representing scenes as compositional generative neural feature fields. In *CVPR*, pages 11453–11464, 2021.

- [38] Maxime Oquab, Timothée Darcet, Theo Moutakanni, Huy V. Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, Russell Howes, Po-Yao Huang, Hu Xu, Vasu Sharma, Shang-Wen Li, Wojciech Galuba, Mike Rabbat, Mido Assran, Nicolas Ballas, Gabriel Synnaeve, Ishan Misra, Herve Jegou, Julien Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski. Dinov2: Learning robust visual features without supervision, 2023.
- [39] Mingjie Pan, Jiaming Liu, Renrui Zhang, Peixiang Huang, Xiaoqi Li, Li Liu, and Shanghang Zhang. Renderocc: Vision-centric 3d occupancy prediction with 2d rendering supervision. *arXiv preprint arXiv:2309.09502*, 2023.
- [40] Ben Poole, Ajay Jain, Jonathan T Barron, and Ben Mildenhall. Dreamfusion: Text-to-3d using 2d diffusion. In *ICLR*, 2022.
- [41] Lingteng Qiu, Guanying Chen, Xiaodong Gu, Qi zuo, Mutian Xu, Yushuang Wu, Weihao Yuan, Zilong Dong, Liefeng Bo, and Xiaoguang Han. Richdreamer: A generalizable normal-depth diffusion model for detail richness in text-to-3d. *arXiv preprint arXiv:2311.16918*, 2023.
- [42] Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. Zero-shot text-to-image generation. In *ICML*, pages 8821–8831. Pmlr, 2021.
- [43] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *CVPR*, pages 10684–10695, 2022.
- [44] Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily L Denton, Kamyar Ghasemipour, Raphael Gontijo Lopes, Burcu Karagol Ayan, Tim Salimans, et al. Photorealistic text-to-image diffusion models with deep language understanding. *NeurIPS*, 35:36479–36494, 2022.
- [45] Johannes L Schonberger and Jan-Michael Frahm. Structure-from-motion revisited. In *CVPR*, pages 4104–4113, 2016.
- [46] Steven M Seitz, Brian Curless, James Diebel, Daniel Scharstein, and Richard Szeliski. A comparison and evaluation of multi-view stereo reconstruction algorithms. In *CVPR*, volume 1, pages 519–528. IEEE, 2006.
- [47] Ruoxi Shi, Hansheng Chen, Zhuoyang Zhang, Minghua Liu, Chao Xu, Xinyue Wei, Linghao Chen, Chong Zeng, and Hao Su. Zero123++: a single image to consistent multi-view diffusion base model. *arXiv preprint arXiv:2310.15110*, 2023.
- [48] Yichun Shi, Peng Wang, Jianglong Ye, Long Mai, Kejie Li, and Xiao Yang. Mvdream: Multi-view diffusion for 3d generation. In *ICLR*, 2023.
- [49] Jaehyeok Shim, Changwoo Kang, and Kyungdon Joo. Diffusion-based signed distance fields for 3d shape generation. In *CVPR*, pages 20887–20897, 2023.
- [50] Vincent Sitzmann, Semon Rezchikov, Bill Freeman, Josh Tenenbaum, and Fredo Durand. Light field networks: Neural scene representations with single-evaluation rendering. *NeurIPS*, 34:19313–19325, 2021.
- [51] Ivan Skorokhodov, Sergey Tulyakov, Yiqun Wang, and Peter Wonka. Epigraf: Rethinking training of 3d gans. *NeurIPS*, 35:24487–24501, 2022.
- [52] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- [53] Cheng Sun, Min Sun, and Hwann-Tzong Chen. Direct voxel grid optimization: Super-fast convergence for radiance fields reconstruction. 2022 ieee. In *CVPR*, pages 5449–5459, 2021.
- [54] Jiaxiang Tang, Zhaoxi Chen, Xiaokang Chen, Tengfei Wang, Gang Zeng, and Ziwei Liu. Lgm: Large multi-view gaussian model for high-resolution 3d content creation. *arXiv preprint arXiv:2402.05054*, 2024.
- [55] Jiaxiang Tang, Jiawei Ren, Hang Zhou, Ziwei Liu, and Gang Zeng. Dreamgaussian: Generative gaussian splatting for efficient 3d content creation. In *ICLR*, 2024.
- [56] Dmitry Tochilkin, David Pankratz, Zexiang Liu, Zixuan Huang, Adam Letts, Yangguang Li, Ding Liang, Christian Laforte, Varun Jampani, and Yan-Pei Cao. Triposr: Fast 3d object reconstruction from a single image. *arXiv preprint arXiv:2403.02151*, 2024.
- [57] Matias Turkulainen, Xuqian Ren, Iaroslav Melekhov, Otto Seiskari, Esa Rahtu, and Juho Kannala. Dn-splatter: Depth and normal priors for gaussian splatting and meshing. *arXiv preprint arXiv:2403.17822*, 2024.

- [58] Shimon Ullman. The interpretation of structure from motion. *Proceedings of the Royal Society of London. Series B. Biological Sciences*, 203(1153):405–426, 1979.
- [59] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *NeurIPS*, 30, 2017.
- [60] Vikram Voleti, Chun-Han Yao, Mark Boss, Adam Letts, David Pankratz, Dmitry Tochilkin, Christian Laforte, Robin Rombach, and Varun Jampani. Sv3d: Novel multi-view synthesis and 3d generation from a single image using latent video diffusion. *arXiv preprint arXiv:2403.12008*, 2024.
- [61] Haochen Wang, Xiaodan Du, Jiahao Li, Raymond A Yeh, and Greg Shakhnarovich. Score jacobian chaining: Lifting pretrained 2d diffusion models for 3d generation. In *CVPR*, pages 12619–12629, 2023.
- [62] Peng Wang, Hao Tan, Sai Bi, Yinghao Xu, Fujun Luan, Kalyan Sunkavalli, Wenping Wang, Zexiang Xu, and Kai Zhang. Pf-lrm: Pose-free large reconstruction model for joint pose and shape prediction. In *ICLR*, 2023.
- [63] Zhengyi Wang, Cheng Lu, Yikai Wang, Fan Bao, Chongxuan Li, Hang Su, and Jun Zhu. Prolificdreamer: High-fidelity and diverse text-to-3d generation with variational score distillation. *NeurIPS*, 36, 2024.
- [64] Zhengyi Wang, Yikai Wang, Yifei Chen, Chendong Xiang, Shuo Chen, Dajiang Yu, Chongxuan Li, Hang Su, and Jun Zhu. Crm: Single image to 3d textured mesh with convolutional reconstruction model. *arXiv preprint arXiv:2403.05034*, 2024.
- [65] Xinyue Wei, Kai Zhang, Sai Bi, Hao Tan, Fujun Luan, Valentin Deschaintre, Kalyan Sunkavalli, Hao Su, and Zexiang Xu. Meshlrm: Large reconstruction model for high-quality mesh. *arXiv preprint arXiv:2404.12385*, 2024.
- [66] Yi Wei, Linqing Zhao, Wenzhao Zheng, Zheng Zhu, Jie Zhou, and Jiwen Lu. Surroundocc: Multi-camera 3d occupancy prediction for autonomous driving. In *ICCV*, pages 21729–21740, 2023.
- [67] Haohan Weng, Tianyu Yang, Jianan Wang, Yu Li, Tong Zhang, CL Chen, and Lei Zhang. Consistent123: Improve consistency for one image to 3d object synthesis. In *ICLR*, 2024.
- [68] Matthew J Westoby, James Brasington, Niel F Glasser, Michael J Hambrey, and Jennifer M Reynolds. ‘structure-from-motion’ photogrammetry: A low-cost, effective tool for geoscience applications. *Geomorphology*, 179:300–314, 2012.
- [69] Sangmin Woo, Byeongjun Park, Hyojun Go, Jin-Young Kim, and Changick Kim. Harmonyview: Harmonizing consistency and diversity in one-image-to-3d. In *CVPR*, 2024.
- [70] Dejia Xu, Ye Yuan, Morteza Mardani, Sifei Liu, Jiaming Song, Zhangyang Wang, and Arash Vahdat. Agg: Amortized generative 3d gaussians for single image to 3d. *arXiv preprint arXiv:2401.04099*, 2024.
- [71] Jiale Xu, Weihao Cheng, Yiming Gao, Xintao Wang, Shenghua Gao, and Ying Shan. Instantmesh: Efficient 3d mesh generation from a single image with sparse-view large reconstruction models. *arXiv preprint arXiv:2404.07191*, 2024.
- [72] Yinghao Xu, Sida Peng, Ceyuan Yang, Yujun Shen, and Bolei Zhou. 3d-aware image synthesis via learning structural and textural representations. In *CVPR*, pages 18430–18439, 2022.
- [73] Yinghao Xu, Zifan Shi, Wang Yifan, Hansheng Chen, Ceyuan Yang, Sida Peng, Yujun Shen, and Gordon Wetzstein. Grm: Large gaussian reconstruction model for efficient 3d reconstruction and generation. *arXiv preprint arXiv:2403.14621*, 2024.
- [74] Yinghao Xu, Hao Tan, Fujun Luan, Sai Bi, Peng Wang, Jiahao Li, Zifan Shi, Kalyan Sunkavalli, Gordon Wetzstein, Zexiang Xu, et al. Dmv3d: Denoising multi-view diffusion using 3d large reconstruction model. In *ICLR*, 2023.
- [75] Taoran Yi, Jiemin Fang, Guanjuan Wu, Lingxi Xie, Xiaopeng Zhang, Wenyu Liu, Qi Tian, and Xinggang Wang. Gaussiandreamer: Fast generation from text to 3d gaussian splatting with point cloud priors. In *CVPR*, 2024.
- [76] Alex Yu, Sara Fridovich-Keil, Matthew Tancik, Qinhong Chen, Benjamin Recht, and Angjoo Kanazawa. Plenoxels: Radiance fields without neural networks. *arXiv preprint arXiv:2112.05131*, 2(3):6, 2021.
- [77] Xin Yu, Yuan-Chen Guo, Yangguang Li, Ding Liang, Song-Hai Zhang, and Xiaojuan Qi. Text-to-3d with classifier score distillation. In *ICLR*, 2024.

- [78] Biao Zhang and Rico Sennrich. Root mean square layer normalization. *NeurIPS*, 32, 2019.
- [79] Biao Zhang, Jiapeng Tang, Matthias Niessner, and Peter Wonka. 3dshape2vecset: A 3d shape representation for neural fields and generative diffusion models. *TOG*, 42(4):1–16, 2023.
- [80] Chubin Zhang, Juncheng Yan, Yi Wei, Jiaxin Li, Li Liu, Yansong Tang, Yueqi Duan, and Jiwen Lu. Occnerf: Self-supervised multi-camera occupancy prediction with neural radiance fields. *arXiv preprint arXiv:2312.09243*, 2023.
- [81] Kai Zhang, Sai Bi, Hao Tan, Yuanbo Xiangli, Nanxuan Zhao, Kalyan Sunkavalli, and Zexiang Xu. Gs-lrm: Large reconstruction model for 3d gaussian splatting. *arXiv preprint arXiv:2404.19702*, 2024.
- [82] Kai Zhang, Gernot Riegler, Noah Snaveley, and Vladlen Koltun. Nerf++: Analyzing and improving neural radiance fields. *arXiv preprint arXiv:2010.07492*, 2020.
- [83] Richard Zhang, Phillip Isola, Alexei A Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *CVPR*, pages 586–595, 2018.
- [84] Linqi Zhou, Yilun Du, and Jiajun Wu. 3d shape generation and completion through point-voxel diffusion. In *ICCV*, pages 5826–5835, 2021.
- [85] Xizhou Zhu, Weijie Su, Lewei Lu, Bin Li, Xiaogang Wang, and Jifeng Dai. Deformable detr: Deformable transformers for end-to-end object detection. In *ICLR*, 2020.

## A Appendix

### A.1 Mesh Extraction from 3D Gaussians

We follow the mesh extraction pipeline from [54] to extract high quality mesh representations from 3D Gaussians. The result is shown in Figure 6.



Figure 6: Image-to-3D generation with mesh extraction results.

### A.2 Occupancy Ground Truth

Previous studies [39, 80, 19] have investigated the task of vision-centric occupancy prediction. However, these approaches often exhibit significant performance discrepancies when compared to 3D methods. To bridge this gap, we leverage depth maps from the GObjaverse dataset to generate accurate 3D occupancy ground truths. This process begins by transforming each pixel in the depth map, represented as  $\mathbf{p}^i = [u, v, 1]^T$ , into a point in world coordinates. This transformation uses both the intrinsic matrix  $K$  and the extrinsic parameters  $T$ , consisting of a rotation matrix  $R$  and a translation vector  $\mathbf{t}$ , as shown in the equation:

$$\mathbf{p}^w = R(d \cdot K^{-1} \mathbf{p}^i) + \mathbf{t}, \quad (14)$$

where  $d$  denotes the depth at pixel  $\mathbf{p}^i$ . Subsequently, these world coordinates are voxelized to pinpoint occupied voxel centres:

$$V = \left\{ \left\lfloor \frac{P}{\epsilon} \right\rfloor \right\} \cdot \epsilon, \quad (15)$$

where  $P$  includes all points in three-dimensional space,  $V$  represents the voxel centers, and the voxel size  $\epsilon$  is set at  $1/128$ . The voxelization helps in reducing redundancy by removing duplicate entries.

### A.3 More Implementation Details

We illustrate the details of network architecture and training procedure in Table 4. We train both the proposal transformer and the reconstruction transformer for 12 epochs on GObjaverse [41], which takes 0.5 and 2 days respectively on 32 A100 40G. For the proposal transformer, we use a batch size of 2 per GPU and apply mixed-precision training with BF16 data type. For the reconstruction transformer, we use a batch size of 1 per GPU and keep the full precision. We note that the second stage is particularly sensitive to the data type and would fail if using mixed-precision.

Table 4: Implementation details.

Proposal Transformer	Image encoder	DINOv2 (ViT-B/14) + Conv
	# layers	6
	# attention head	16
	# deformed points	8
	Image feature dimension	384
	3D feature dimension	384
	Max sequence length	4096
Reconstruction Transformer	Image encoder	DINOv2 (ViT-B/14) + Conv
	# layers	16
	# attention head	16
	# deformed points	8
	Image feature dimension	384
	3D feature dimension	768
	Max sequence length	4096
Training details	# Gaussians per token	32
	Epoch	12
	Learning rate	1e-4
	Learning rate scheduler	Cosine
	Optimizer	AdamW
	(Beta1, Beta2)	(0.9, 0.95)
	Weight decay	0.05
	Warm-up	1500
	Gradient accumulation	8
	Gradient clip	4
	# GPU	32